

Dilated Neural Networks for Time Series Forecasting

Chenhui Hu

Data Scientist
Microsoft Cloud & Enterprise

Agenda

- Overview of Time Series Forecasting Methods
- Foundations of Dilated Convolutional Neural Networks
- Application to Retail Demand Forecasting

Agenda


- Overview of Time Series Forecasting Methods
- Foundations of Dilated Convolutional Neural Networks
- Application to Retail Demand Forecasting

Overview

Time Series Forecasting Applications



retail sales forecasting

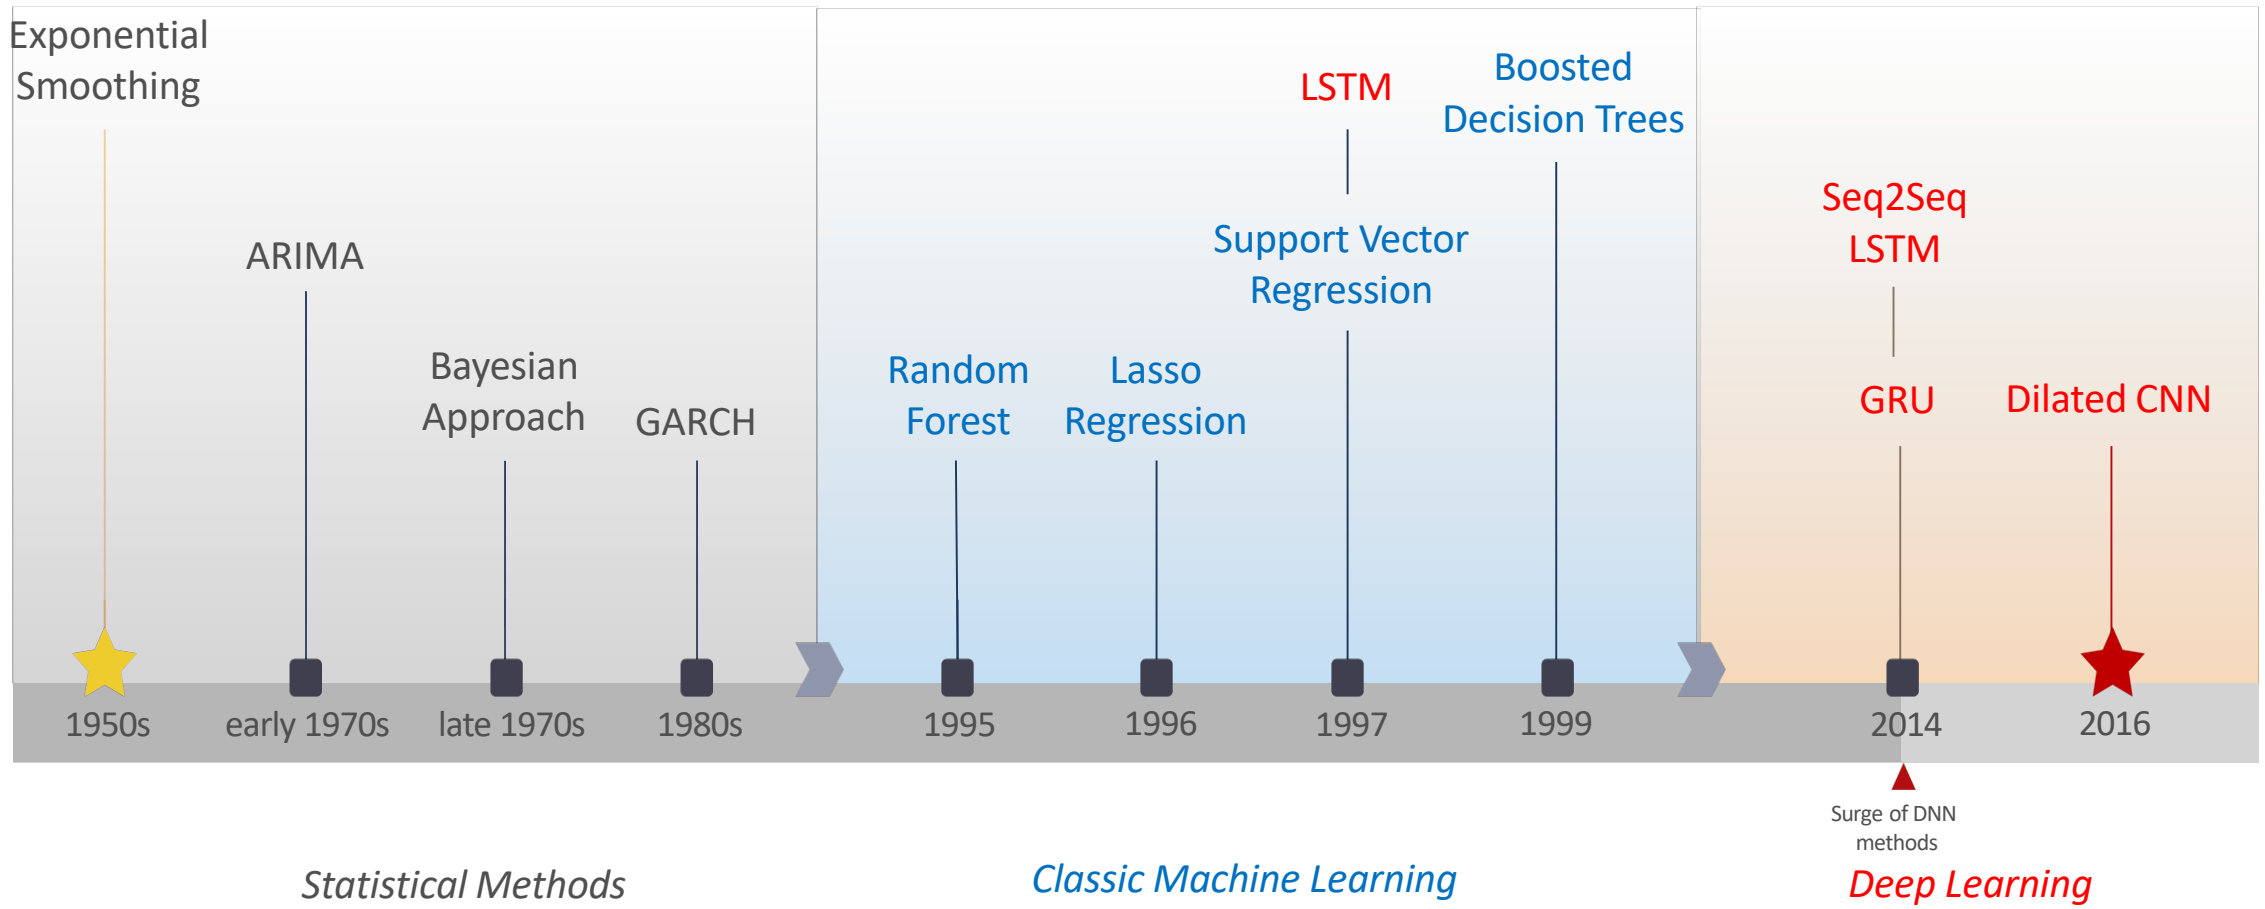


financial forecasting



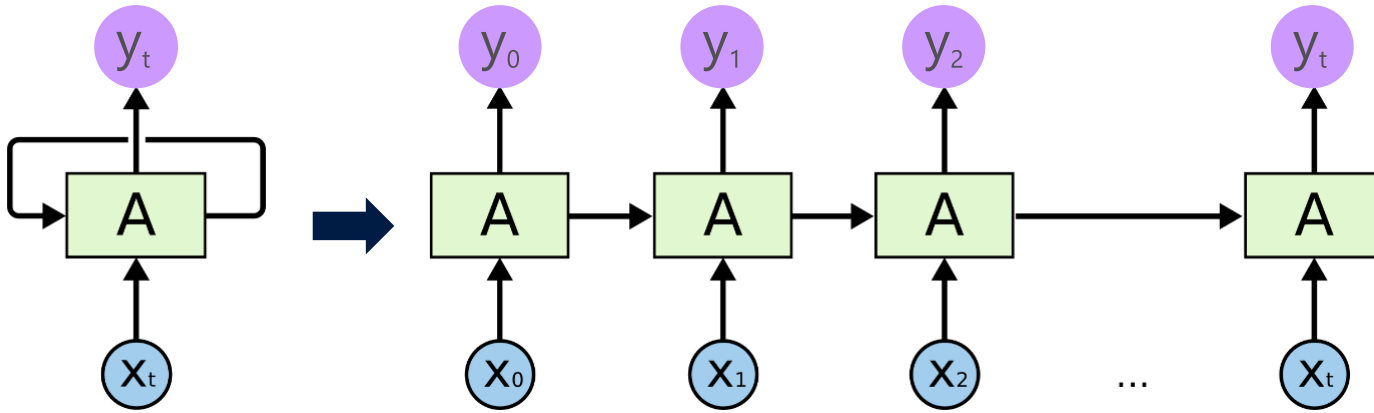
web traffic forecasting

Representative Methods




Recurrent NN

- RNN models *the order* of the data explicitly
- Long Short-Term Memory (LSTM) network models long-range dependencies




unrolled Recurrent Neural Network (RNN)



multi-layer perceptron

Convolutional NN

- Success of CNN in CV and recently in NLP
 - Outperforms human in many computer vision tasks
 - Achieves state-of-the-art accuracy in audio synthesis & machine translation
- MLP with convolution layers



2D convolution as a linear combination of pixel values in a neighborhood

RNN vs. CNN

- RNN
 - Has been default DL choice for sequence modeling
 - Can learn long-range dependencies effectively
 - But training is usually slower compared with CNN
- CNN
 - Lower level of model complexity
 - Easy to parallelize the computation
 - Dilated CNN can outperform RNN in sequence modeling

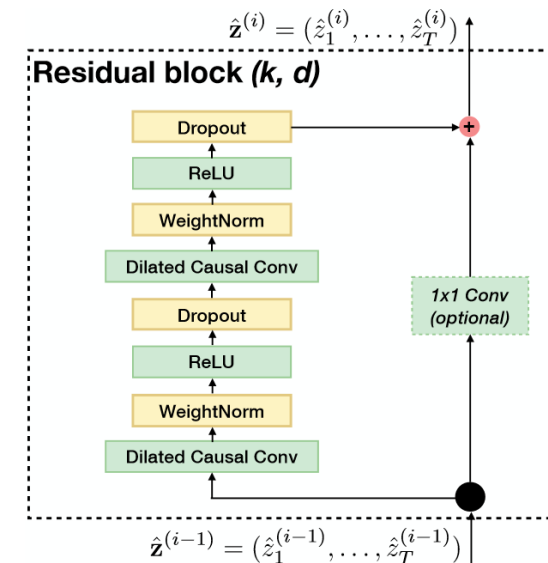



Agenda

- Overview of Time Series Forecasting Methods
- Foundations of Dilated Convolutional Neural Networks
- Application to Retail Demand Forecasting

Dilated CNN Overview

- CNN with dilated causal conv. layers
 - Basic setup: dilated conv. layers + output layer
 - Advanced: skip connections, dropout, batch normalization, gated activation units, ...
- Examples
 - WaveNet for audio synthesis
 - Temporal CNN for sequence modeling



1-D Convolution for Time Series

↓

	Feature 1	Feature 2	Feature 3	Feature 4
Timestamp 1	0.4	0.8	0.5	0.7
Timestamp 2	0.7	0.3	0.6	0.7
Timestamp 3	0.8	0.4	0.9	0.2
Timestamp 4	0.1	0.5	0.8	0.5
Timestamp 5	0.4	0.8	0.7	0.1
Timestamp 6	0.3	0.0	0.2	0.6
Timestamp 7	0.2	0.5	0.8	0.2

Filter

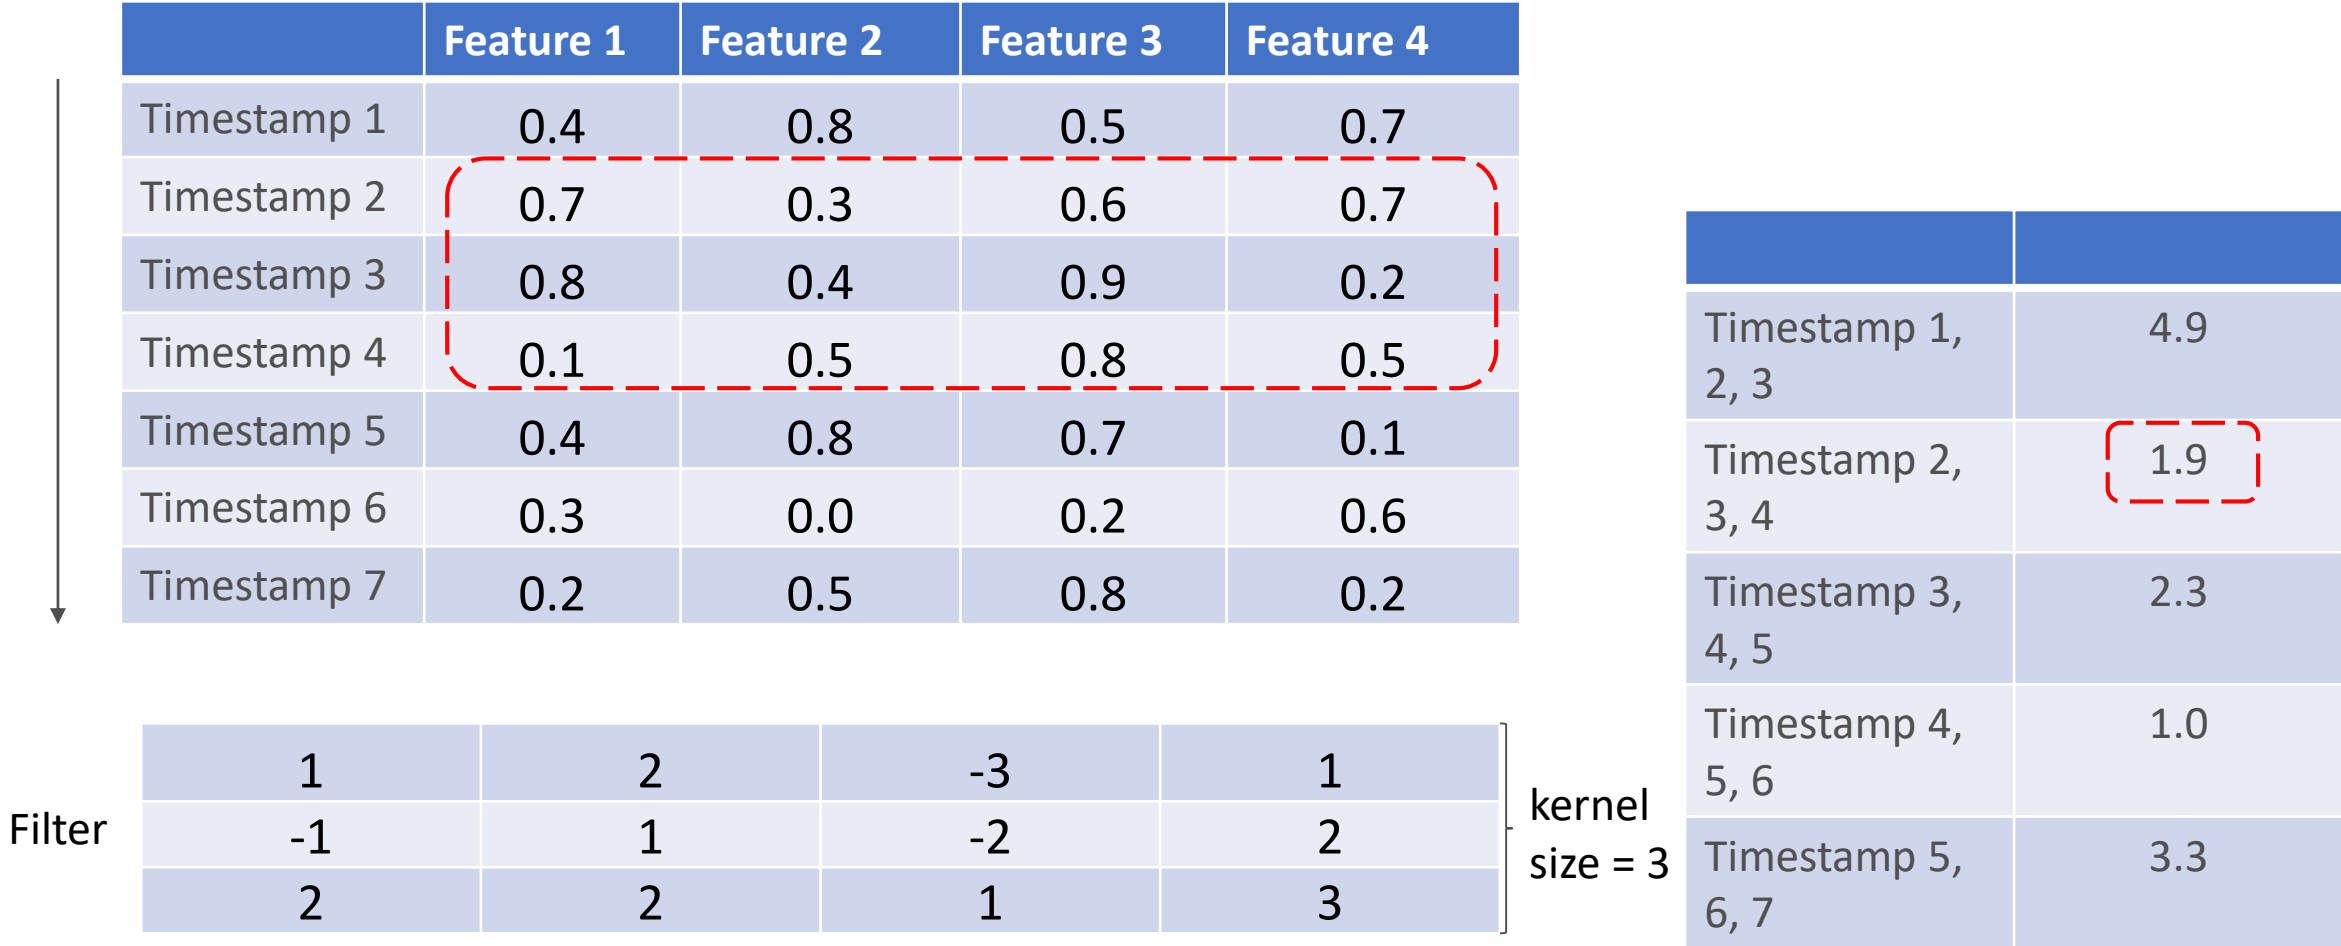
1	2	-3	1
-1	1	-2	2
2	2	1	3

kernel size = 3

$$4.9 = 0.4 \times 1 + 0.8 \times 2 + 0.5 \times (-3) + 0.7 \times 1 + 0.7 \times (-1) + 0.3 \times 1 + 0.6 \times (-2) + 0.7 \times 2 + 0.8 \times 2 + 0.4 \times 2 + 0.9 \times 1 + 0.2 \times 3$$

Timestamp 1, 2, 3	4.9
Timestamp 2, 3, 4	1.9
Timestamp 3, 4, 5	2.3
Timestamp 4, 5, 6	1.0
Timestamp 5, 6, 7	3.3

1-D Convolution for Time Series



1-D Convolution with Multiple Filters

	Feature 1	Feature 2	Feature 3	Feature 4
Timestamp 1	0.4	0.8	0.5	0.7
Timestamp 2	0.7	0.3	0.6	0.7
Timestamp 3	0.8	0.4	0.9	0.2
Timestamp 4	0.1	0.5	0.8	0.5
Timestamp 5	0.4	0.8	0.7	0.1
Timestamp 6	0.3	0.0	0.2	0.6
Timestamp 7	0.2	0.5	0.8	0.2

Filters

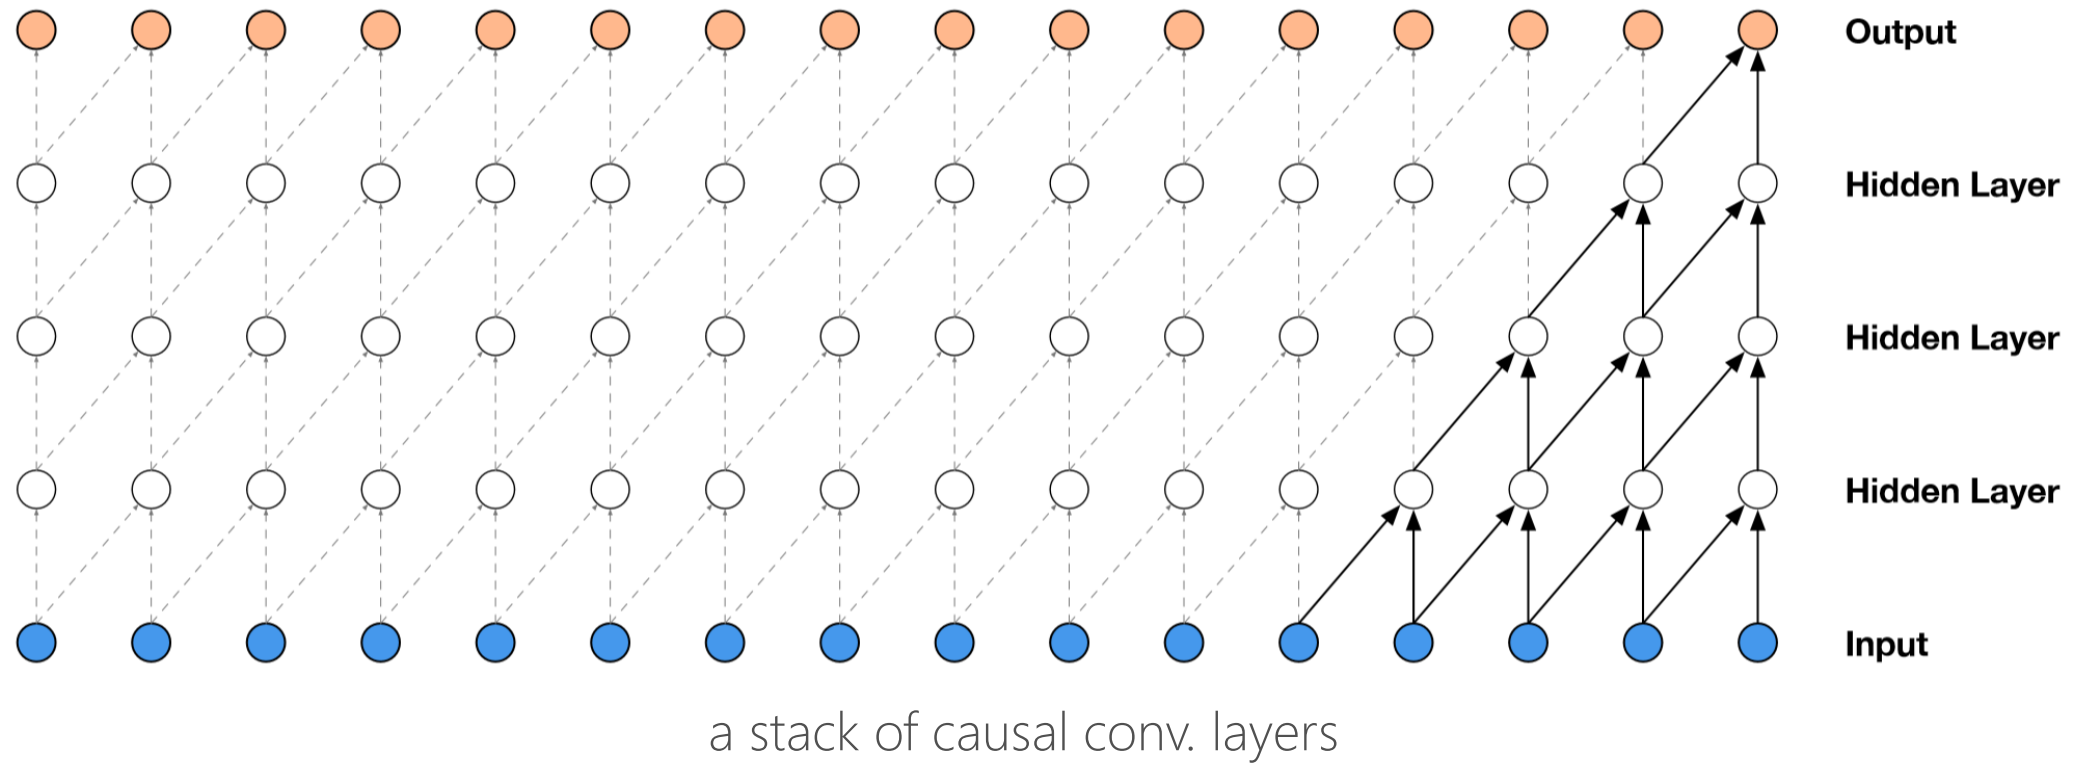
1	2	-3	1
-1	1	-2	2
2	2	1	3

1	1	2	1
1	0	1	0
1	1	2	1

	Filter 1	Filter 2
Timestamp 1, 2, 3	4.9	7.4
Timestamp 2, 3, 4	1.9	7.3
Timestamp 3, 4, 5	2.3	6.8
Timestamp 4, 5, 6	1.0	5.1
Timestamp 5, 6, 7	3.3	5.7

Causal Convolution

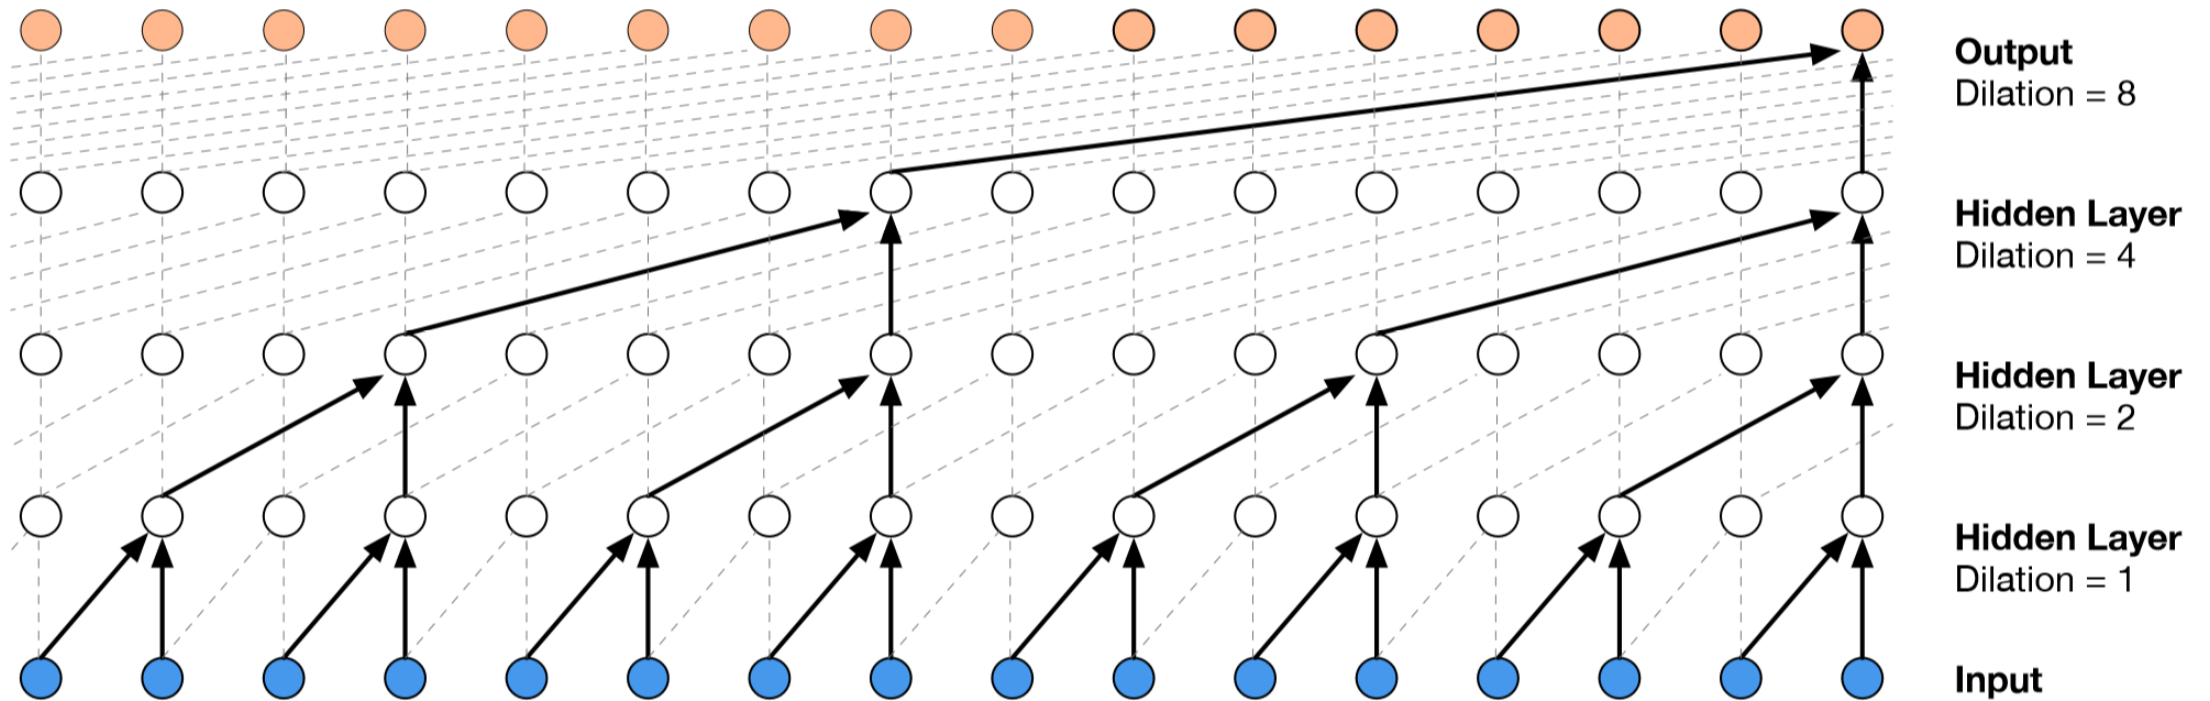
Only existing (no future) data is used in each neuron



Dilated Convolution

With dilation width d , the conv. window starts at location i of size k is

$$\left[\mathbf{x}_i \quad \mathbf{x}_{i+d} \quad \mathbf{x}_{i+2d} \quad \cdots \quad \mathbf{x}_{i+(k-1) \cdot d} \right]$$



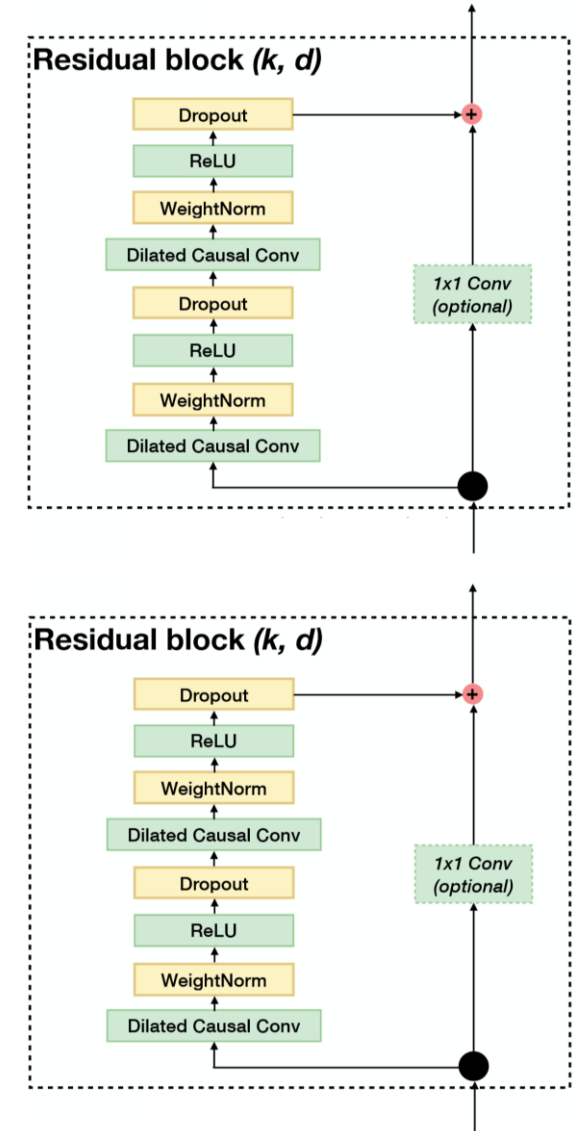
stack of *dilated* causal convolutional layers

Dilated-CNN Advantages


- Increases receptive field
 - Receptive field is the part of the data visible to a neuron
 - Exponentially increasing the dilation factor results in exponential receptive field growth
- Captures global view of the input with less parameters
- Handles temporal flow with causal structure

Dilated-CNN Design

- ReLU activation: $h = \max(0, a)$, where $a = Wx + b$.
- Dropout for regularization
- Use WeightNorm or BatchNorm to speed up convergence
- Use skip connection to avoid degradation of model performance when adding more layers
- Concatenate building blocks to increase model capacity



Dilated-RNNs



example of a three-layer Dilated-RNN with dilation 1, 2, and 4

Agenda

- Overview of Time Series Forecasting Methods
- Foundations of Dilated Convolutional Neural Networks
- Application to Retail Demand Forecasting

Application to Retail Demand Forecasting

OrangeJuice dataset in *bayesm* R package

store	brand	week	logmove	constant	price1	price2	price3
2	1	40	9.018695	1	0.060469	0.060497	0.042031
2	1	46	8.723231	1	0.060469	0.060312	0.045156
2	1	47	8.253228	1	0.060469	0.060312	0.045156
2	1	48	8.987197	1	0.060469	0.060312	0.049844


price4	price5	price6	price7	price8	price9	price10
0.029531	0.049531	0.053021	0.038906	0.041406	0.028906	0.024844
0.046719	0.049531	0.047813	0.045781	0.027969	0.042969	0.042031
0.046719	0.037344	0.053021	0.045781	0.041406	0.048125	0.032656
0.037344	0.049531	0.053021	0.045781	0.041406	0.042344	0.032656

price11	deal	feat	profit
0.038984	1	0.0	37.992326
0.038984	0	0.0	30.126667
0.038984	0	0.0	30.000000
0.038984	0	0.0	29.950000

83 stores
11 brands
121 weeks

Sample Data

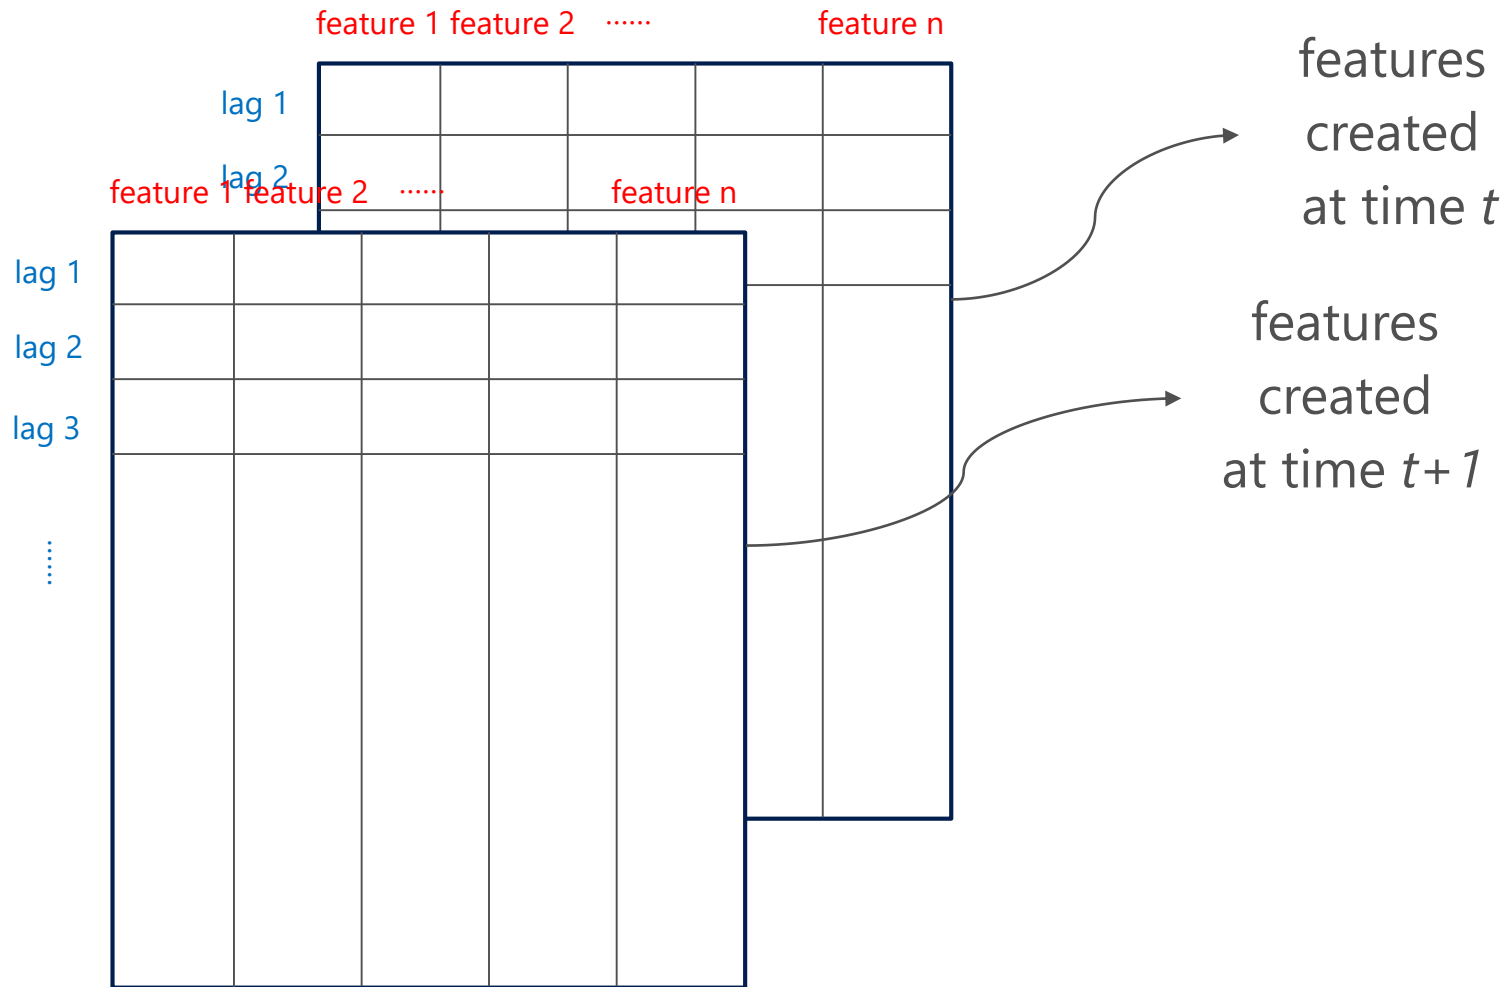
Weekly sales of store 2 brand 1 (missing values are filled with zeros)



Feature Engineering

- *move* – historical unit sales
- *deal, feat* – if on sale or advertisement
- *price* – price of the current brand
- *price_ratio* – relative price to other brands
$$\frac{\textit{price}}{\textit{average_price_of_all_brands}}$$
- *month* – month number
- *week_of_month* – week number of the month

Data Transformation

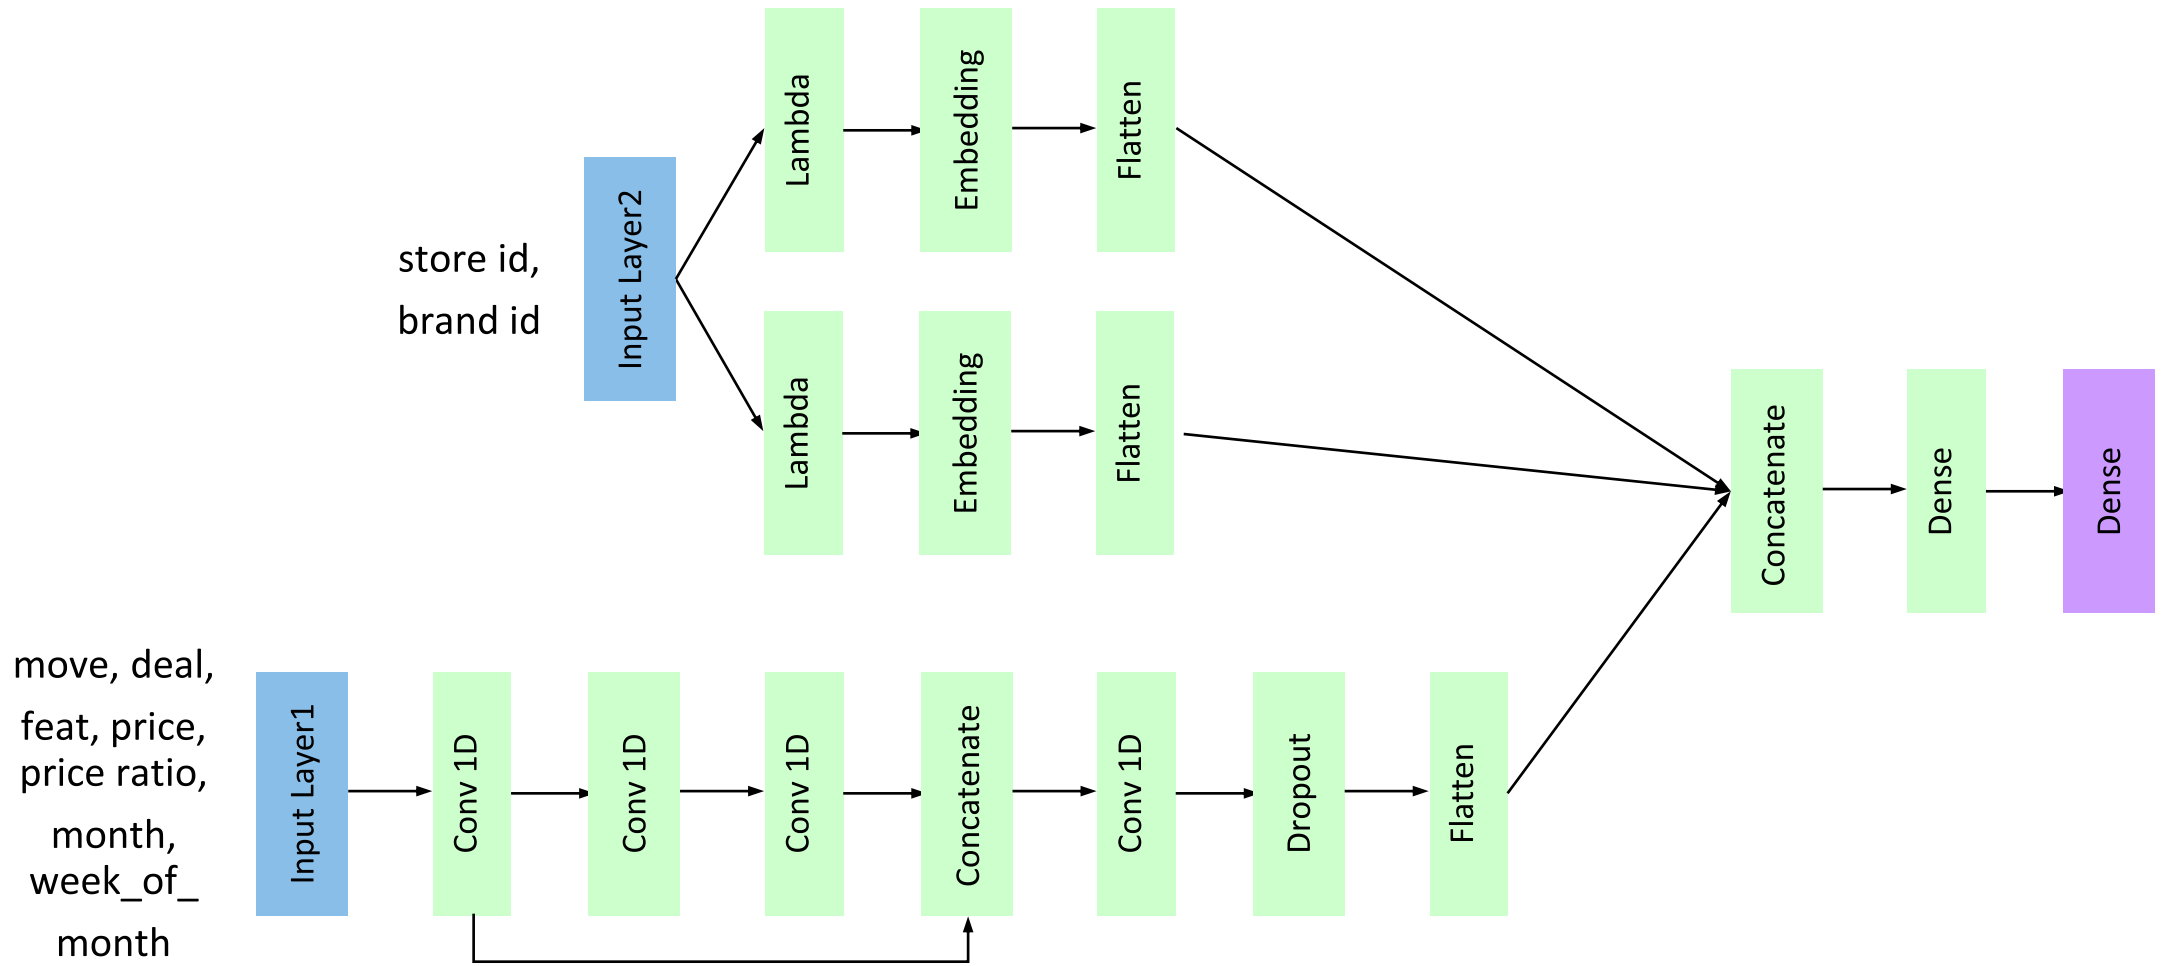


(# of lags, # of features)



(# of samples,
of lags, # of features)

Network Structure



[#samples, 15, 7] ➡ [#samples, 15, 3]

Model Definition

- *Embedding*: encode categorical features into binary sequences
- *Conv1D*
 - n_filters
 - kernel_size
 - dilation_rate
 - padding
 - activation function
- *Skip connections*
- *Combine with categorial features and pass to a dense layer*

```
# Sequential input
seq_in = Input(shape=(seq_len, n_input_series))

# Categorical input
cat_fea_in = Input(shape=(2,), dtype='uint8')
store_id = Lambda(lambda x: x[:, 0, None])(cat_fea_in)
brand_id = Lambda(lambda x: x[:, 1, None])(cat_fea_in)
store_embed = Embedding(MAX_STORE_ID+1, 7, input_length=1)(store_id)
brand_embed = Embedding(MAX_BRAND_ID+1, 4, input_length=1)(brand_id)

# Dilated convolutional layers
c1 = Conv1D(filters=n_filters, kernel_size=kernel_size, dilation_rate=1,
            padding='causal', activation='relu')(seq_in)
c2 = Conv1D(filters=n_filters, kernel_size=kernel_size, dilation_rate=2,
            padding='causal', activation='relu')(c1)
c3 = Conv1D(filters=n_filters, kernel_size=kernel_size, dilation_rate=4,
            padding='causal', activation='relu')(c2)

# Skip connections
c4 = concatenate([c1, c3])

# Output of convolutional layers
conv_out = Conv1D(8, 1, activation='relu')(c4)
conv_out = Dropout(args.dropout_rate)(conv_out)
conv_out = Flatten()(conv_out)

# Concatenate with categorical features
x = concatenate([conv_out, Flatten()(store_embed), Flatten()(brand_embed)])
x = Dense(16, activation='relu')(x)
output = Dense(n_outputs, activation='linear')(x)

# Define model interface, loss function, and optimizer
model = Model(inputs=[seq_in, cat_fea_in], outputs=output)
```

Model Definition

- *Embedding*: encode categorical features into binary sequences
- *Conv1D*
 - n_filters
 - kernel_size
 - dilation_rate
 - padding
 - activation function
- *Skip connections*
- *Combine with categorial features and pass to a dense layer*

```
# Sequential input
seq_in = Input(shape=(seq_len, n_input_series))

# Categorical input
cat_fea_in = Input(shape=(2,), dtype='uint8')
store_id = Lambda(lambda x: x[:, 0, None])(cat_fea_in)
brand_id = Lambda(lambda x: x[:, 1, None])(cat_fea_in)
store_embed = Embedding(MAX_STORE_ID+1, 7, input_length=1)(store_id)
brand_embed = Embedding(MAX_BRAND_ID+1, 4, input_length=1)(brand_id)

# Dilated convolutional layers
c1 = Conv1D(filters=n_filters, kernel_size=kernel_size, dilation_rate=1,
            padding='causal', activation='relu')(seq_in)
c2 = Conv1D(filters=n_filters, kernel_size=kernel_size, dilation_rate=2,
            padding='causal', activation='relu')(c1)
c3 = Conv1D(filters=n_filters, kernel_size=kernel_size, dilation_rate=4,
            padding='causal', activation='relu')(c2)

# Skip connections
c4 = concatenate([c1, c3])

# Output of convolutional layers
conv_out = Conv1D(8, 1, activation='relu')(c4)
conv_out = Dropout(args.dropout_rate)(conv_out)
conv_out = Flatten()(conv_out)

# Concatenate with categorical features
x = concatenate([conv_out, Flatten()(store_embed), Flatten()(brand_embed)])
x = Dense(16, activation='relu')(x)
output = Dense(n_outputs, activation='linear')(x)

# Define model interface, loss function, and optimizer
model = Model(inputs=[seq_in, cat_fea_in], outputs=output)
```

Model Definition

- *Embedding*: encode categorical features into binary sequences
- *Conv1D*
 - n_filters
 - kernel_size
 - dilation_rate
 - padding
 - activation function
- *Skip connections*
- *Combine with categorical features and pass to a dense layer*

```
# Sequential input
seq_in = Input(shape=(seq_len, n_input_series))

# Categorical input
cat_fea_in = Input(shape=(2,), dtype='uint8')
store_id = Lambda(lambda x: x[:, 0, None])(cat_fea_in)
brand_id = Lambda(lambda x: x[:, 1, None])(cat_fea_in)
store_embed = Embedding(MAX_STORE_ID+1, 7, input_length=1)(store_id)
brand_embed = Embedding(MAX_BRAND_ID+1, 4, input_length=1)(brand_id)

# Dilated convolutional layers
c1 = Conv1D(filters=n_filters, kernel_size=kernel_size, dilation_rate=1,
            padding='causal', activation='relu')(seq_in)
c2 = Conv1D(filters=n_filters, kernel_size=kernel_size, dilation_rate=2,
            padding='causal', activation='relu')(c1)
c3 = Conv1D(filters=n_filters, kernel_size=kernel_size, dilation_rate=4,
            padding='causal', activation='relu')(c2)

# Skip connections
c4 = concatenate([c1, c3])

# Output of convolutional layers
conv_out = Conv1D(8, 1, activation='relu')(c4)
conv_out = Dropout(args.dropout_rate)(conv_out)
conv_out = Flatten()(conv_out)

# Concatenate with categorical features
x = concatenate([conv_out, Flatten()(store_embed), Flatten()(brand_embed)])
x = Dense(16, activation='relu')(x)
output = Dense(n_outputs, activation='linear')(x)

# Define model interface, loss function, and optimizer
model = Model(inputs=[seq_in, cat_fea_in], outputs=output)
```

Model Definition

- *Embedding*: encode categorical features into binary sequences
- *Conv1D*
 - n_filters
 - kernel_size
 - dilation_rate
 - padding
 - activation function
- *Skip connections*
- *Combine with categorical features and pass to a dense layer*

```
# Sequential input
seq_in = Input(shape=(seq_len, n_input_series))

# Categorical input
cat_fea_in = Input(shape=(2,), dtype='uint8')
store_id = Lambda(lambda x: x[:, 0, None])(cat_fea_in)
brand_id = Lambda(lambda x: x[:, 1, None])(cat_fea_in)
store_embed = Embedding(MAX_STORE_ID+1, 7, input_length=1)(store_id)
brand_embed = Embedding(MAX_BRAND_ID+1, 4, input_length=1)(brand_id)

# Dilated convolutional layers
c1 = Conv1D(filters=n_filters, kernel_size=kernel_size, dilation_rate=1,
            padding='causal', activation='relu')(seq_in)
c2 = Conv1D(filters=n_filters, kernel_size=kernel_size, dilation_rate=2,
            padding='causal', activation='relu')(c1)
c3 = Conv1D(filters=n_filters, kernel_size=kernel_size, dilation_rate=4,
            padding='causal', activation='relu')(c2)

# Skip connections
c4 = concatenate([c1, c3])

# Output of convolutional layers
conv_out = Conv1D(8, 1, activation='relu')(c4)
conv_out = Dropout(args.dropout_rate)(conv_out)
conv_out = Flatten()(conv_out)

# Concatenate with categorical features
x = concatenate([conv_out, Flatten()(store_embed), Flatten()(brand_embed)])
x = Dense(16, activation='relu')(x)
output = Dense(n_outputs, activation='linear')(x)

# Define model interface, loss function, and optimizer
model = Model(inputs=[seq_in, cat_fea_in], outputs=output)
```

Hyperparameter Tuning

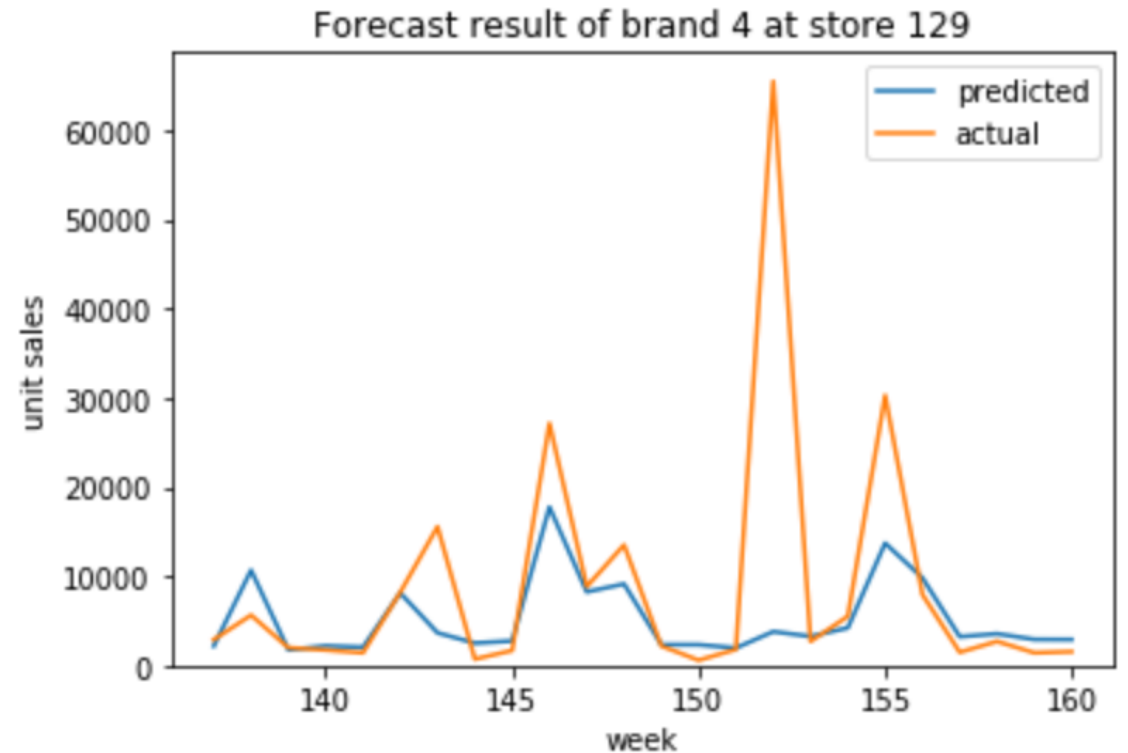
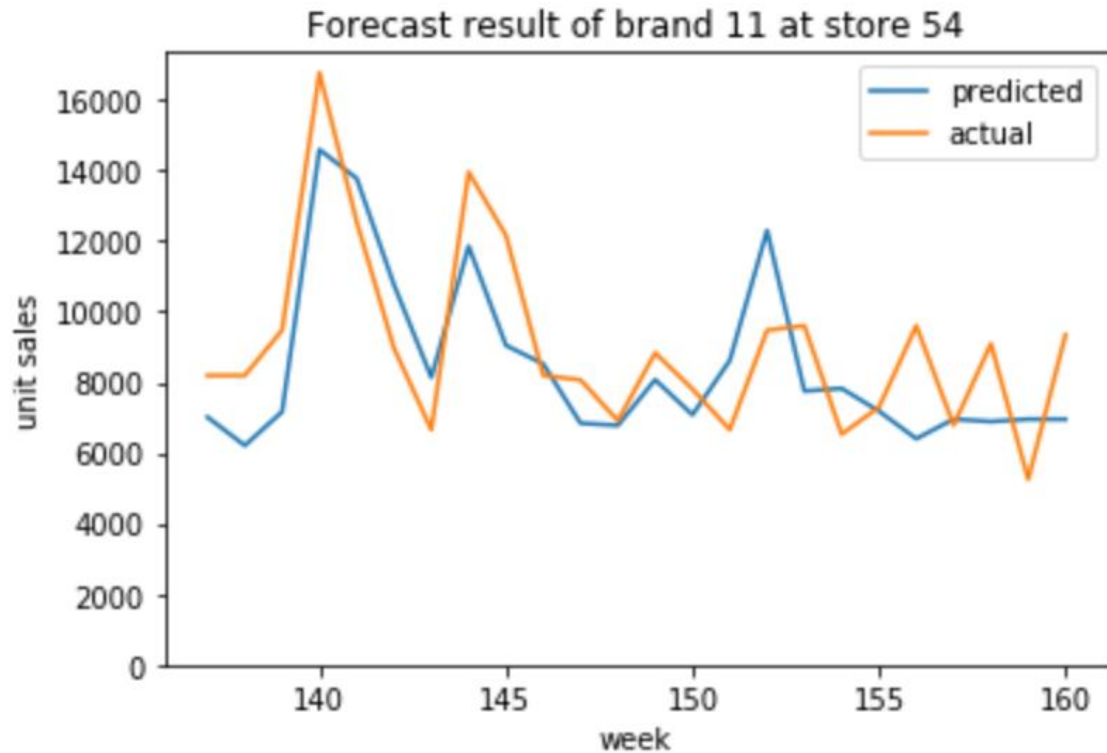
HyperDrive via Azure Machine Learning SDK

- Automates the hyperparameter sweeps on an elastic cluster
- Supports Bayesian sampling and early termination

```
ps = BayesianParameterSampling({
    '--seq-len': quniform(5, 40, 1),
    '--dropout-rate': uniform(0, 0.4),
    '--batch-size': choice(32, 64),
    '--learning-rate': choice(1e-4, 1e-3, 5e-3, 1e-2, 1.5e-2, 2e-2, 3e-2, 5e-2, 1e-1),
    '--epochs': quniform(2, 80, 1)
})
htc = HyperDriveRunConfig(estimator=est,
                          hyperparameter_sampling=ps,
                          primary_metric_name='MAPE',
                          primary_metric_goal=PrimaryMetricGoal.MINIMIZE,
                          max_total_runs=200,
                          max_concurrent_runs=4)
htr = exp.submit(config=htc)
```

Results

- Forecast unit sales between week 138 and 160 with model retrained every 2 weeks
- Azure Ubuntu Linux VM with one-half Tesla K80 GPU, 56 GB memory



Results

Mean absolute percentage error (MAPE)

Method	MAPE	Running time	Machine
Dilated CNN	37.09 %	413 s	GPU Linux VM
Seq2Seq RNN	37.68 %	669 s	GPU Linux VM
Naive	109.67 %	114.06 s	CPU Linux VM
ETS	70.99 %	277.01 s	CPU Linux VM
ARIMA	70.80 %	265.94 s	CPU Linux VM

Results are collected based on the median of 5 run results

Thank you!

Rate today's session

Cyberconflict: A new era of war, sabotage, and fear

David Sanger (The New York Times)
9:55am-10:10am Wednesday, March 27, 2019
Location: Ballroom
Secondary topics: Security and Privacy

Rate This Session

See passes & pricing

Add to Your Schedule
Add Comment or Question

We're living in a new era of constant sabotage, misinformation, and fear, in which everyone is a target, and you're often the collateral damage in a growing conflict among states. From crippling infrastructure to sowing discord and doubt, cyber is now the weapon of choice for democracies, dictators, and terrorists.

David Sanger explains how the rise of cyberweapons has transformed geopolitics like nothing since the invention of the atomic bomb. Moving from the White House Situation Room to the dens of Chinese, Russian, North Korean, and Iranian hackers to the boardrooms of Silicon Valley, David reveals a world coming face-to-face with the perils of technological revolution—a conflict that the United States helped start when it began using cyberweapons against Iranian nuclear plants and North Korean missile launches. But now we find ourselves in a conflict we're uncertain how to control, as our adversaries exploit vulnerabilities in our hyperconnected nation and we struggle to figure out how to deter these complex, short-of-war attacks.

David Sanger
The New York Times

David E. Sanger is the national security correspondent for the *New York Times* as well as a national security and political contributor for CNN and a frequent guest on *CBS This Morning*, *Face the Nation*, and many PBS shows.




Session page on conference website

✓ Attending Notes Remove

Cyberconflict: A new era of war, sabotage, and fear

🕒 9:55 AM - 10:10 AM, Wed, Mar 27, 2019

Speakers



David Sanger
National Security Correspondent
The New York Times

📍 Ballroom

Keynotes

David Sanger explains how the rise of cyberweapons has transformed geopolitics like nothing since the invention of the atomic bomb. From crippling infrastructure to sowing discord and doubt, cyber is now the weapon of choice for democracies, dictators, and terrorists.

SESSION EVALUATION

O'Reilly Events App