




DEEP LEARNING IN PYTHON


# **The need for optimization**

# A baseline neural network



- Actual Value of Target: 13
- Error: Predicted - Actual = -4

# A baseline neural network



- Actual Value of Target: 13
- Error: Predicted - Actual = 0

# Predictions with multiple points

- Making accurate predictions gets harder with more points
- At any set of weights, there are many values of the error
- ... corresponding to the many points we make predictions for



# Loss function

- Aggregates errors in predictions from many data points into single number
- Measure of model's predictive performance




# Squared error loss function

| Prediction | Actual | Error | Squared Error |
|------------|--------|-------|---------------|
| 10         | 20     | -10   | 100           |
| 8          | 3      | 5     | 25            |
| 6          | 1      | 5     | 25            |

- Total Squared Error: 150
- Mean Squared Error: 50

# Loss function



Loss function





# Loss function

- Lower loss function value means a better model
- Goal: Find the weights that give the lowest value for the loss function
- Gradient descent




# Gradient descent

- Imagine you are in a pitch dark field
- Want to find the lowest point
- Feel the ground to see how it slopes
- Take a small step downhill
- Repeat until it is uphill in every direction

# Gradient descent steps

- Start at random point
- Until you are somewhere flat:
  - Find the slope
  - Take a step downhill

# Optimizing a model with a single weight





DEEP LEARNING IN PYTHON


**Let's practice!**



DEEP LEARNING IN PYTHON

# Gradient descent

# Gradient descent





# Gradient descent

- If the slope is positive:
  - Going opposite the slope means moving to lower numbers
  - Subtract the slope from the current value
  - Too big a step might lead us astray
- Solution: learning rate
  - Update each weight by subtracting **learning rate \* slope**





# Slope calculation example



- To calculate the slope for a weight, need to multiply:
  - Slope of the loss function w.r.t value at the node we feed into
  - The value of the node that feeds into our weight
  - Slope of the activation function w.r.t value we feed into



# Slope calculation example



- To calculate the slope for a weight, need to multiply:
  - Slope of the loss function w.r.t value at the node we feed into
  - The value of the node that feeds into our weight
  - Slope of the activation function w.r.t value we feed into



# Slope calculation example



- Slope of mean-squared loss function w.r.t prediction:
  - $2 * (\text{Predicted Value} - \text{Actual Value}) = 2 * \text{Error}$
  - $2 * -4$



# Slope calculation example



- To calculate the slope for a weight, need to multiply:
  - Slope of the loss function w.r.t value at the node we feed into
  - The value of the node that feeds into our weight
  - Slope of the activation function w.r.t value we feed into



# Slope calculation example



- To calculate the slope for a weight, need to multiply:
  - Slope of the loss function w.r.t value at the node we feed into
  - The value of the node that feeds into our weight
  - Slope of the activation function w.r.t value we feed into



# Slope calculation example



- To calculate the slope for a weight, need to multiply:
  - Slope of the loss function w.r.t value at the node we feed into
  - The value of the node that feeds into our weight
  - ~~Slope of the activation function w.r.t value we feed into~~



# Slope calculation example




- $2 * -4 * 3$
- -24
- If learning rate is 0.01, the new weight would be
- $2 - 0.01(-24) = 2.24$





# Network with two inputs affecting prediction





# Code to calculate slopes and update weights

```
In [1]: import numpy as np

In [2]: weights = np.array([1, 2])

In [3]: input_data = np.array([3, 4])

In [4]: target = 6

In [5]: learning_rate = 0.01

In [6]: preds = (weights * input_data).sum()

In [7]: error = preds - target

In [8]: print(error)
5
```



# Code to calculate slopes and update weights

```
In [9]: gradient = 2 * input_data * error
```

```
In [10]: gradient
```

```
Out[10]: array([30, 40])
```

```
In [11]: weights_updated = weights - learning_rate * gradient
```

```
In [12]: preds_updated = (weights_updated * input_data).sum()
```

```
In [13]: error_updated = preds_updated - target
```

```
In [14]: print(error_updated)
```

```
-2.5
```



DEEP LEARNING IN PYTHON

**Let's practice!**




DEEP LEARNING IN PYTHON

# Backpropagation



# Backpropagation

- Allows gradient descent to update all weights in neural network (by getting gradients for all weights)
- Comes from chain rule of calculus
- Important to understand the process, but you will generally use a library that implements this



# Backpropagation process


- Trying to estimate the slope of the loss function w.r.t each weight
- Do forward propagation to calculate predictions and errors






# Backpropagation process

ReLU Activation Function  
Actual Target Value = 4



ReLU Activation Function  
Actual Target Value = 4  
Error = 3






# Backpropagation process

- Go back one layer at a time
- Gradients for weight is product of:
  1. Node value feeding into that weight
  2. Slope of loss function w.r.t node it feeds into
  3. Slope of activation function at the node it feeds into



# ReLU Activation Function



# Backpropagation process

- Need to also keep track of the slopes of the loss function w.r.t node values
- Slope of node values are the sum of the slopes for all weights that come out of them



DEEP LEARNING IN PYTHON

**Let's practice!**




DEEP LEARNING IN PYTHON

# Backpropagation in practice





# Backpropagation




ReLU Activation Function  
Actual Target Value = 4  
Error = 3

- Top weight's slope =  $1 * 6$
- Bottom weight's slope =  $3 * 6$



# Backpropagation






# Calculating slopes associated with any weight

- Gradients for weight is product of:
  1. Node value feeding into that weight
  2. Slope of activation function for the node being fed into
  3. Slope of loss function w.r.t output node



# Backpropagation



| Current Weight Value | Gradient |
|----------------------|----------|
| 0                    | 0        |
| 1                    | 6        |
| 2                    | 0        |
| 3                    | 18       |

Lower level slope of weight is current output node's level slope of loss, it multiply incoming node value to get slope of weight for the current weight.



# Backpropagation: Recap

- Start at some random set of weights
- Use forward propagation to make a prediction
- Use backward propagation to calculate the slope of the loss function w.r.t each weight
- Multiply that slope by the learning rate, and subtract from the current weights
- Keep going with that cycle until we get to a flat part

# Stochastic gradient descent

- It is common to calculate slopes on only a subset of the data ('batch')
- Use a different batch of data to calculate the next update
- Start over from the beginning once all data is used
- Each time through the training data is called an epoch
- When slopes are calculated on one batch at a time: stochastic gradient descent



DEEP LEARNING IN PYTHON

**Let's practice**